

ISIT312 Big Data Management

# HDFS Interfaces

Dr Fenghui Ren

School of Computing and Information Technology -  
University of Wollongong

# HDFS Interfaces

## Outline

Hadoop Cluster vs. Pseudo-Distributed Hadoop

Shell Interface to HDFS

Web Interface to HDFS

Java Interface to HDFS

Internals of HDFS

# Hadoop Cluster vs. Pseudo-Distributed Hadoop

A **Hadoop** cluster is deployed in a cluster of computer nodes

- As **Hadoop** is developed in Java, all **Hadoop** services sit on Java Virtual Machines running on the cluster nodes

**Hadoop** provides a pseudo-distributed mode on a single machine

- All Java Virtual Machines for necessary **Hadoop** services are running on a single machine
- In our case this machine is a Virtual Machine running under Ubuntu 14.04

**HDFS** provides the following interfaces to read, write, interrogate, and manage the filesystem

- The file system shell (Command-Line Interface): **hadoop fs** or **hdfs dfs**
- **Hadoop** Filesystem Java API
- **Hadoop** simple Web User Interface
- Other interfaces, such as RESTful proxy interfaces (e.g., HttpFS)

# HDFS Interfaces

## Outline

[Hadoop Cluster vs. Pseudo-Distributed Hadoop](#)

[Shell Interface to HDFS](#)

[Web Interface to HDFS](#)

[Java Interface to HDFS](#)

[Internals of HDFS](#)

# Shell Interface to HDFS

Commands are provided in the shell Bash

```
$ which bash  
/bin/bash
```

Bash shell

Hadoop's home directory

```
$ cd $HADOOP_HOME  
$ ls  
bin include libexec logs README.txt share  
etc lib LICENSE.txt NOTICE.txt sbin
```

Home of Hadoop

You will mostly use scripts in the **bin** and **sbin** folders, and use jar files in the **share** folder

Hadoop Daemons

```
$ jps  
28530 SecondaryNameNode  
11188 NodeManager  
28133 NameNode  
28311 DataNode  
10845 ResourceManager  
3542 Jps
```

Hadoop daemons

Hadoop is running properly only if the above services are running

# Shell Interface to HDFS

Create a **HDFS** user account (already created in a virtual machine used by us)

```
$ bin/hadoop fs -mkdir -p /user/bigdata
```

Creating home of user account

Create an folder **input**

```
$ bin/hadoop fs -mkdir input
```

Creating a folder

View the folders in **Hadoop** home

```
$ bin/hadoop fs -ls
```

Found 1 item

```
drwxr-xr-x - bigdata supergroup 0 2017-07-17 16:33 input
```

Listing home of user account

Upload a file to **HDFS**

```
$ bin/hadoop fs -put README.txt input
```

```
$ bin/hadoop fs -ls input
```

```
-rw-r--r-- 1 bigdata supergroup 1494 2017-07-12 17:53 input/README.txt
```

Uploading a file

Read a file in **HDFS**

```
$ bin/hadoop fs -cat input/README.txt
```

<contents of README.txt goes here>

Listing a file

# Shell Interface to HDFS

The path in **HDFS** is represented as a URI with the prefix **hdfs://**

For example

- **hdfs://<hostname>:<port>/user/bigdata/input** refers to the **input** directory in **HDFS** under the user of **bigdata**
- **hdfs ://<hostname>:<port>/user/bigdata/input/README.txt** refers to the file **README.txt** in the above **input** directory in **HDFS**

When interacting with **HDFS** interface in the default setting, one can omit IP, port, and user, and simply mention the directory or file

Thus, the full-spelling of **hadoop fs -ls input** is

```
hadoop fs -ls hdfs://<hostname>:  
<port>/user/bigdata/input
```

# Shell Interface to HDFS

Some of frequently used commands

Commands of Hadoop shell interface

Command	Description
-put	Upload a file (or files) from the local filesystem to HDFS
-mkdir	Create a directory in HDFS
-ls	List the files in a directory in HDFS
-cat	Read the content of a file (or files) in HDFS
-copyFromLocal	Copy a file from the local filesystem to HDFS (similar to put)
-copyToLocal	Copy a file (or files) from HDFS to the local filesystem
-rm	Delete a file (or files) in HDFS
-rm -r	Delete a directory in HDFS



# HDFS Interfaces

## Outline

Hadoop Cluster vs. Pseudo-Distributed Hadoop

Shell Interface to HDFS

Web Interface to HDFS

Java Interface to HDFS

Internals of HDFS

# Web Interface of HDFS

## Overview 'localhost:8020' (active)

<b>Started:</b>	Thu Jul 27 10:53:21 AEST 2017
<b>Version:</b>	2.7.3, rbaa91f7c6bc9cb92be5982de4719c1c8af91ccff
<b>Compiled:</b>	2016-08-18T01:41Z by root from branch-2.7.3
<b>Cluster ID:</b>	CID-50dd23ff-1ec0-4860-b1fd-f8b0067f5b57
<b>Block Pool ID:</b>	BP-680435313-10.0.2.15-1499005796923

## Summary

Security is off.

Safemode is off.

126 files and directories, 62 blocks = 188 total filesystem object(s).

Heap Memory used 39.84 MB of 263 MB Heap Memory. Max Heap Memory is 889 MB.

Non Heap Memory used 46.08 MB of 46.84 MB Committed Non Heap Memory. Max Non Heap Memory

# Web Interface of HDFS

## Browse Directory

/							Go!
Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	bigdata	supergroup	0 B	12/07/2017, 1:42:53 pm	0	0 B	<a href="#">hbase_data</a>
drwxrwx---	bigdata	supergroup	0 B	10/07/2017, 2:23:00 pm	0	0 B	<a href="#">tmp</a>
drwxr-xr-x	bigdata	supergroup	0 B	03/07/2017, 1:29:24 am	0	0 B	<a href="#">user</a>

Hadoop, 2016.

# HDFS Interfaces

## Outline

[Hadoop Cluster vs. Pseudo-Distributed Hadoop](#)

[Shell Interface to HDFS](#)

[Web Interface to HDFS](#)

[Java Interface to HDFS](#)

[Internals of HDFS](#)

# Java Interface to HDFS

A file in a **Hadoop** filesystem is represented by a **Hadoop Path** object

- Its syntax is URI
- For example,  
`hdfs://localhost:8020/user/bigdata/input/README.txt`

To get an instance of **FileSystem**, use the following factory methods

```
public static FileSystem get(Configuration conf) throws IOException
public static FileSystem get(URI uri, Configuration conf) throws IOException
public static FileSystem get(URI uri, Configuration conf, String user)
                                throws IOException
```

Factory methods

The following method gets a local filesystem instance

```
public static FileSystem getLocal(Configuration conf) throws IOException
```

Get local file system method

# Java interface to HDFS

A `Configuration` object is determined by the `Hadoop` configuration files or user-provided parameters

Using the default configuration, one can simply set

```
Configuration conf = new Configuration()
```

Configuration object

With a `FileSystem` instance in hand, we invoke an `open()` method to get the input stream for a file

```
public FSDataInputStream open(Path f) throws IOException  
public abstract FSDataInputStream open(Path f, int bufferSize) throws IOException
```

Open method

A `Path` object can be created by using a designated URI

```
Path f = new Path(uri)
```

Path object

# Java interface to HDFS

Putting together, we can create the following file reading application

```
public class FileSystemCat {  
    public static void main(String[] args) throws Exception {  
        String uri = args[0];  
        Configuration conf = new Configuration();  
        FileSystem fs = FileSystem.get(URI.create(uri), conf);  
        FSDataInputStream in = null;  
        Path path = new Path(uri);  
        in = fs.open(path);  
        IOUtils.copyBytes(in, System.out, 4096, true);  
    }  
}
```

Class FileSystemCat

# Java interface to HDFS

The compilation simply uses the `javac` command, but it needs to point the dependencies in the class path.

```
export HADOOP_CLASSPATH=$(HADOOP_HOME/bin/hadoop classpath)
javac -cp $HADOOP_CLASSPATH FileSystemCat.java
```

Compilation

Then, a `jar` file is created and run as follows

```
jar cvf FileSystemCat.jar FileSystemCat*.class
hadoop jar FileSystemCat.jar FileSystemcat input/README.txt
```

jar file and processing

The output is the same as processing a command `hadoop fs -cat`



# Java interface to HDFS

Suppose an input stream is created to read a local file

To write a file on **HDFS**, the simplest way is to take a **Path** object for the file to be created and return an output stream to write to

```
public FSDataOutputStream create(Path f) throws IOException
```

Path object

And then just copy the input stream to the output stream

Another, more flexible, way is to read the input stream into a buffer and then write to the output stream

# Java interface to HDFS

## A file writing application

File writing

```
public class FileSystemPut {  
    public static void main(String[] args) throws Exception {  
        String localStr = args[0];  
        String hdfsStr = args[1];  
        Configuration conf = new Configuration();  
        FileSystem local = FileSystem.getLocal(conf);  
        FileSystem hdfs = FileSystem.get(URI.create(hdfsStr), conf);  
        Path localFile = new Path(localStr);  
        Path hdfsFile = new Path(hdfsStr);  
        FSDataInputStream in = local.open(localFile);  
        FSDataOutputStream out = hdfs.create(hdfsFile);  
        IOUtils.copyBytes(in, out, 4096, true);  
    }  
}
```

# Java interface to HDFS

## Another file writing application

```
public class FileSystemPutAlt {  
    public static void main(String[] args) throws Exception {  
        String localStr = args[0];  
        String hdfsStr = args[1];  
        Configuration conf = new Configuration();  
        FileSystem local = FileSystem.getLocal(conf);  
        FileSystem hdfs = FileSystem.get(URI.create(hdfsStr), conf);  
        Path localFile = new Path(localStr);  
        Path hdfsFile = new Path(hdfsStr);  
        FSDataInputStream in = local.open(localFile);  
        FSDataOutputStream out = hdfs.create(hdfsFile);  
        byte[] buffer = new byte[256];  
        int bytesRead = 0;  
        while( (bytesRead = in.read(buffer)) > 0 ) {  
            out.write(buffer, 0, bytesRead);  
        }  
        in.close();  
        out.close();  
    }  
}
```

File writing

# Java interface to HDFS

Other file system API methods

The method `mkdirs()` creates a directory

The method `getFileStatus()` gets the meta information for a single file or directory

The method `listStatus()` lists contents of files in a directory

The method `exists()` checks whether a file exists

The method `delete()` removes a file

The Java API enables the implementation of customised applications to interact with [HDFS](#)

# HDFS Interfaces

## Outline

[Hadoop Cluster vs. Pseudo-Distributed Hadoop](#)

[Shell Interface to HDFS](#)

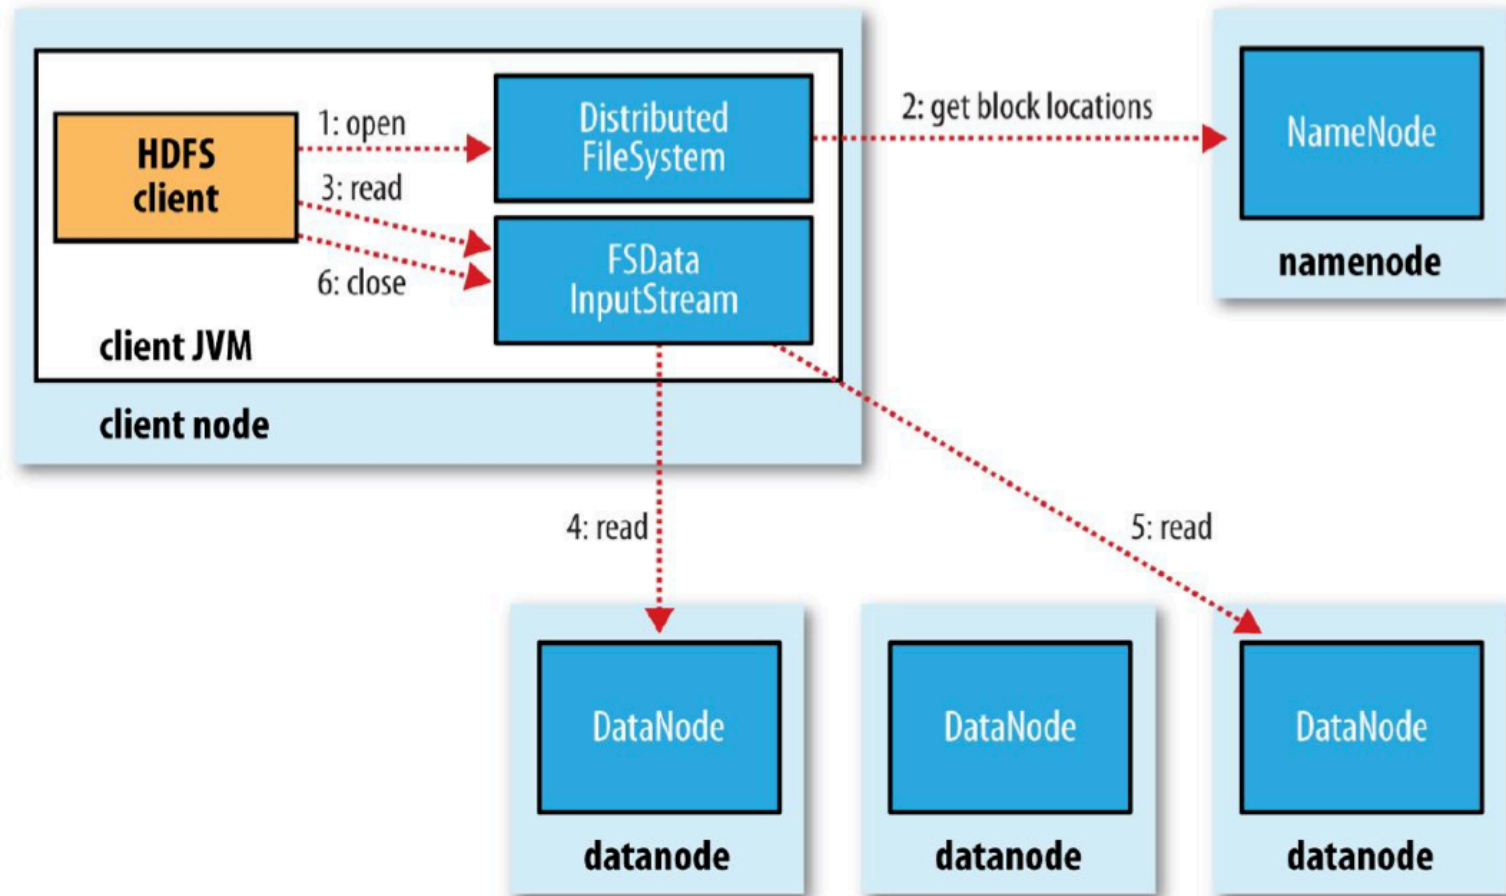
[Web Interface to HDFS](#)

[Java Interface to HDFS](#)

[Internals of HDFS](#)

# Internals of HDFS

What happens "inside" when we read data into HDFS ?



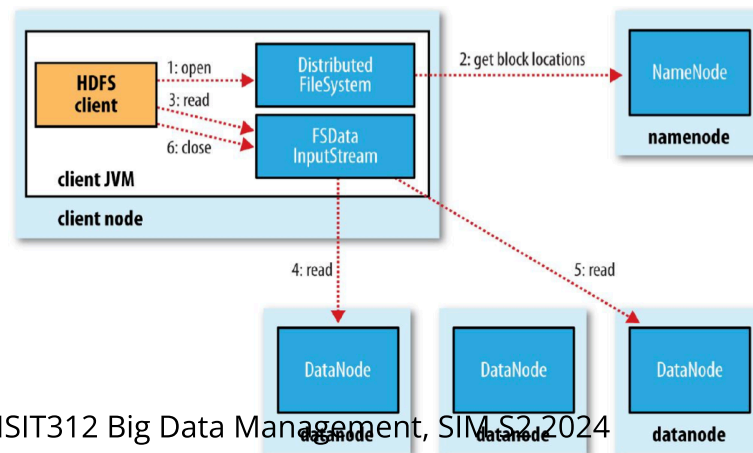
# Internals of HDFS

Read data from HDFS

Step 1: The client opens the file it wishes to read by calling `open()` on the `FileSystem` object, which for HDFS is an instance of `DistributedFileSystem`

Step 2: `DistributedFileSystem` calls the `namenode`, using remote procedure calls (RPCs), to determine the locations of the first few blocks in the file

Step 3: The `DistributedFileSystem` returns an `FSDaataInputStream` to the client and the client calls `read()` on the stream

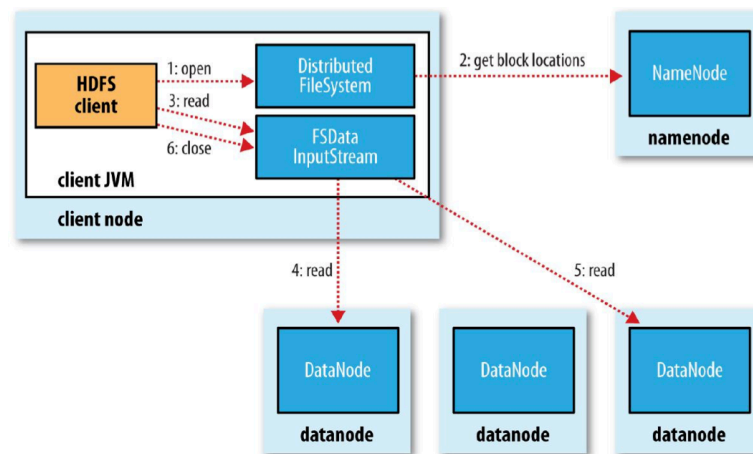


# Internals of HDFS

Step 4: `FSDaataInputStream` connects to the first datanode for the first block in the file, and then data is streamed from the datanode back to the client, by calling `read()` repeatedly on the stream

Step 5: When the end of the block is reached, `FSDaataInputStream` will close the connection to the datanode, then find the best (possibly the same) datanode for the next block

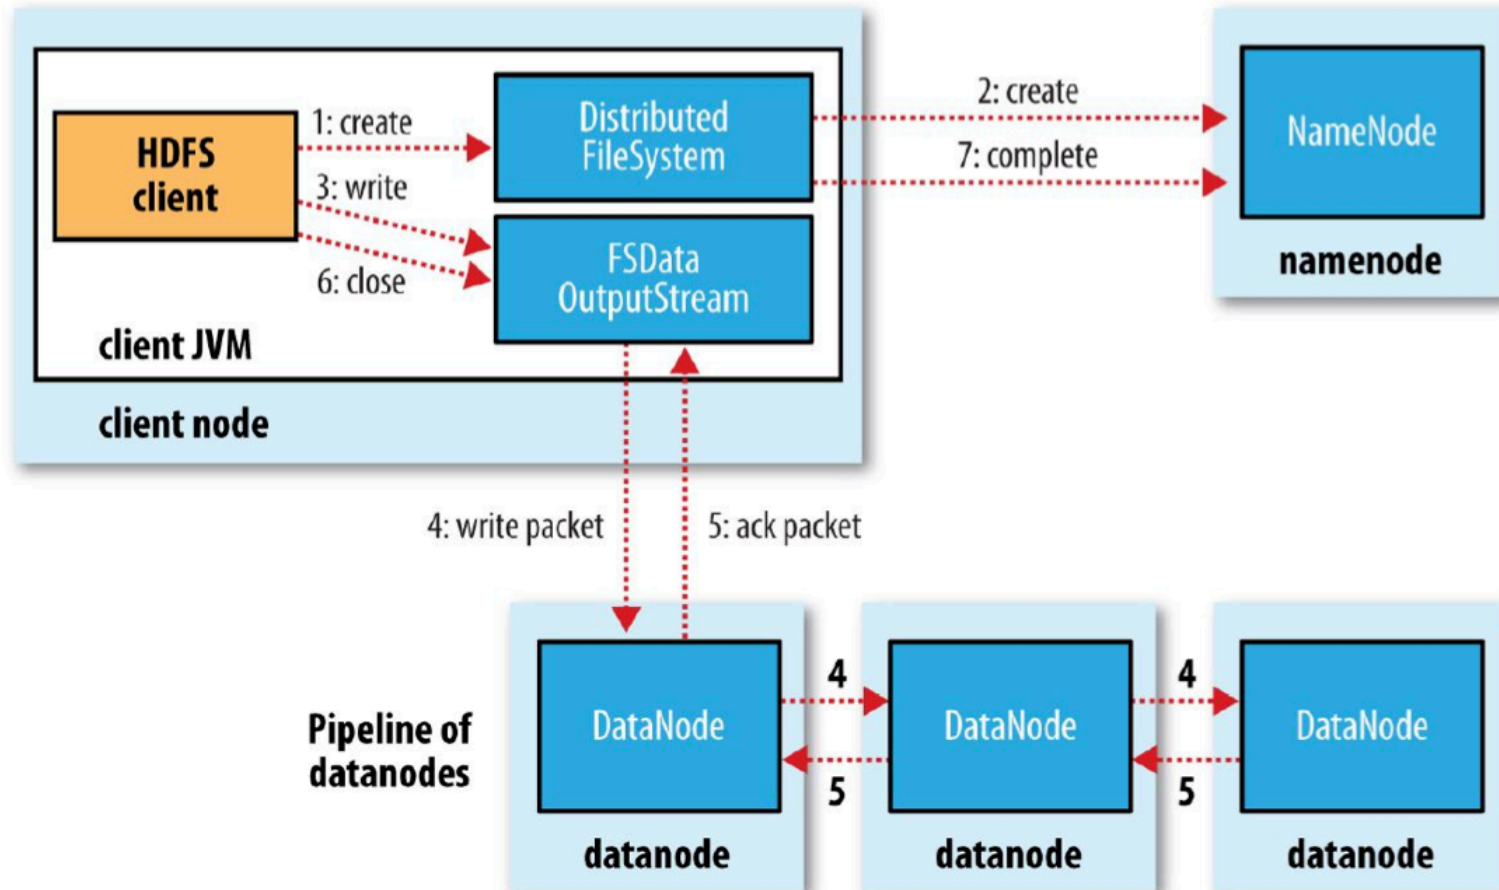
Step 6: When the client has finished reading, it calls `close()` on the `FSDaataInputStream`





# Internals of HDFS

Write data into HDFS



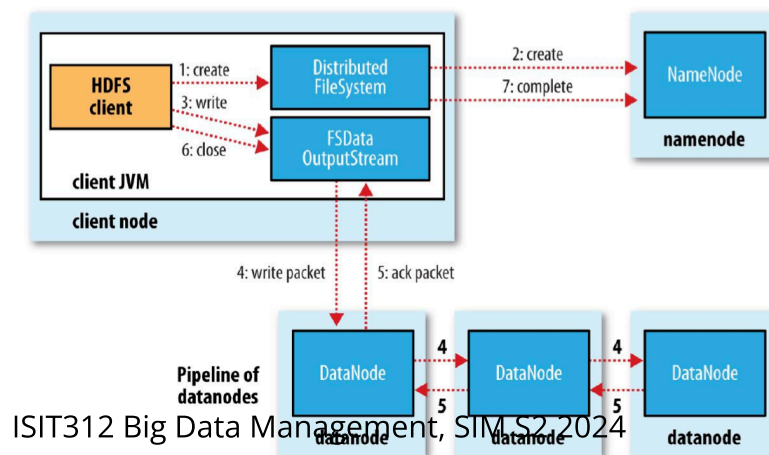
# Internals of HDFS

Step 1: The client creates the file by calling `create()` on `DistributedFileSystem`

Step 2: `DistributedFileSystem` makes an RPC call to the `namenode` to create a new file in the file system namespace and returns an `FSDaataOutputStream` for the client to start writing data to

Step 3: The client writes data into the `FSDaataOutputStream`

Step 4: Data wrapped by the `FSDaataOutputStream` is split into packages, which are flushed into a queue; data packages are sent to the blocks in a datanode and forwarded to other (usually two) datanodes

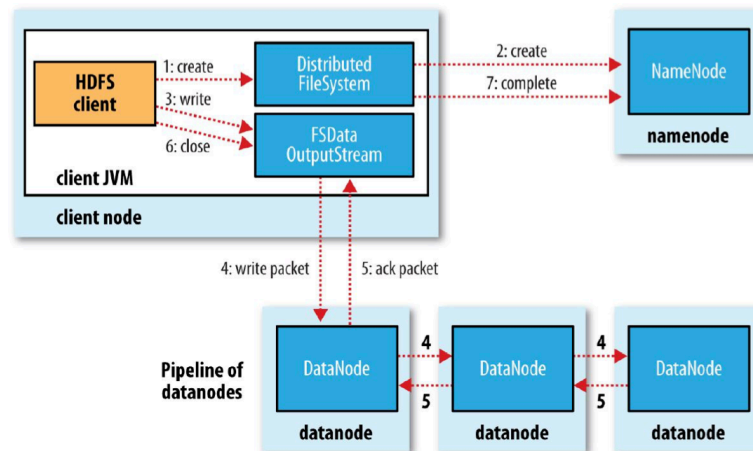


# Internals of HDFS

Step 5: If `FSDataStream` receives an ack signal from the datanode the data packages are removed from the queue

Step 6: When the client has finished writing data, it calls `close()` on the stream

Step 7: The client signals the namenode that the writing is completed



# References

Vohra D., Practical Hadoop ecosystem: a definitive guide to Hadoop-related frameworks and tools, Apress, 2016 (Available through UOW library)

Aven J., Hadoop in 24 Hours, SAMS Teach Yourself, SAMS 2017