

ISIT312 Big Data Management

# Hive Data Structures

Dr Fenghui Ren

School of Computing and Information Technology -  
University of Wollongong

# Hive Data Structures

## Outline

Primitive Data Types

Complex Data Types

Databases

Tables

Partitions

Buckets

Views

# Primitive Data Types

**TINYINT**, 1 byte, example: 10Y

**SMALLINT**, 2 bytes, example: 10S

**INT**, 4 bytes, example: 10

**BIGINT**, 8 bytes, example: 10L

**FLOAT**, 4 bytes, example: 0.1234567

**DOUBLE**, 8 bytes, example: 0.1234567891234

**DECIMAL**, (m,n), example: 3.14

**BINARY**, n bytes, example: 1011001

**BOOLEAN**, 1 byte example: TRUE

**STRING**, 2G bytes, example: 'Abcdef '

**CHAR**, 255 bytes, example: 'Hello '

# Primitive Data Types

`VARCHAR`, 1 byte, example: 'Hive'

`DATE`, YYYY-MM-DD, example: '2017-05-03'

`TIMESTAMP`, YYYY-MM-DD HH:MM:SS[.fff...] example: '2017-05-03  
15:10:00.345'

# Hive Data Structures

## Outline

[Primitive Data Types](#)

[Complex Data Types](#)

[Databases](#)

[Tables](#)

[Partitions](#)

[Buckets](#)

[Views](#)

# Complex Data Types

**ARRAY:** list of values of the same types,

```
example: ['Hadoop', 'Pig', 'Hive']  
access: bigdata[1]
```

Array type

**MAP:** a set of key-value pairs,

```
example: {'k1': 'Hadoop', 'k2': 'Pig'}  
access: bigdata['k2']
```

Map type

**STRUCT:** user defined structure of any type of fields,

```
example: {name: 'Hadoop', age: 24, salary: 50000.06}  
access: bigdata.name
```

Struct type

# Complex Data Types

The following **CREATE TABLE** command creates a table **types** with complex data types columns

```
CREATE TABLE types(  
  array_col array<string>,  
  map_col map<int,string>,  
  struct_col struct<a:string, b:int, c:double> );
```

A table with columns of types

**INSERT INTO** and **SELECT** statements can be used to load data into a table with complex data types columns

```
INSERT INTO types  
  SELECT array('bolt', 'nut', 'screw'),  
         map(1,'bolt', 2,'nut', 3,'screw'),  
         named_struct('a','bolt', 'b',5, 'c',0.5)  
FROM DUAL;
```

Inserting values into a table

# Hive Data Structures

## Outline

[Primitive Data Types](#)

[Complex Data Types](#)

[Databases](#)

[Tables](#)

[Partitions](#)

[Buckets](#)

[Views](#)



# Databases

**Database** is a collection of conceptually related tables, i.e. tables that implement a conceptual schema

**Database** is implemented as a folder/directory in **HDFS**

A **default database** is located at `/user/hive/warehouse`

A new database is created in a folder `/user/hive/warehouse`

For example, a database `tpchr` is located at `/user/hive/warehouse/tpchr.db`

# Databases

The following **CREATE DATABASE** command creates a database **tpchr**

```
CREATE DATABASE tpchr;
```

Creating a database

To find more information about a database we can use **DESCRIBE DATABASE** command

```
DESCRIBE DATABASE tpchr;
```

Listing a database

A command **USE** makes a database "current" (there is no need to prefix a table name with a database name)

```
USE tpchr;
```

Making a database current

To delete a database we can use **DROP DATABASE** command

```
DROP DATABASE tpchr;
```

Dropping a database

# Hive Data Structures

## Outline

[Primitive Data Types](#)

[Complex Data Types](#)

[Databases](#)

[Tables](#)

[Partitions](#)

[Buckets](#)

[Views](#)

# Tables

An **internal table** (or **managed table**) is a table created by **Hive** in **HDFS**

If data is already stored in **HDFS** then an **external Hive table** can be created to provide a tabular view of the data

Location in **HDFS** of data stored in an external table is specified in the **LOCATION** properties instead of the default warehouse directory

**Hive** fully manages the life cycle (add/delete data, create/drop table) of **internal tables** and data in the **internal tables**

When an **external table** is deleted its metadata information is deleted from a **metastore** and the data is kept in **HDFS**

# Tables

CREATE TABLE statement creates an **internal table**

Creating an internal table

```
CREATE TABLE IF NOT EXISTS intregion(  
  R_REGIONKEY DECIMAL(12),  
  R_NAME VARCHAR(25),  
  R_COMMENT VARCHAR(152) )  
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '|'   
  STORED AS TEXTFILE;
```

LOAD DATA statement loads data into an **internal table**

Loading data into an internal table

```
LOAD DATA LOCAL INPATH 'region.tbl' INTO TABLE intregion;
```

# Tables

CREATE EXTERNAL TABLE statement creates an external table

Creating an external table

```
CREATE EXTERNAL TABLE IF NOT EXISTS extregion(  
  R_REGIONKEY DECIMAL(12),  
  R_NAME VARCHAR(25),  
  R_COMMENT VARCHAR(152) )  
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '|'   
  STORED AS TEXTFILE LOCATION '/user/tpchr/region';
```

LOAD DATA statement loads data into an external table

Loading data into an external table

```
LOAD DATA LOCAL INPATH 'region.tbl' INTO TABLE extregion;
```

# Tables

An **external table** can be created "over" an already existing file in **HDFS**

Loading a file to HDFS

```
hadoop fs -mkdir /user/tpchr/nation
hadoop fs -put nation.tbl /user/tpchr/nation
hadoop fs -ls /user/tpchr/nation
-rw-r--r-- 3 janusz supergroup 401 2017-07-02 10:24 /user/tpchr/nation/nation.tbl
```

Creating an external table over a file in HDFS

```
CREATE EXTERNAL TABLE IF NOT EXISTS extnation(
  N_NATIONKEY DECIMAL(12),
  N_NAME      VARCHAR(25),
  N_COMMENT   VARCHAR(152) )
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '|'
  STORED AS TEXTFILE LOCATION '/user/tpchr/nation';
```

# Hive Data Structures

## Outline

[Primitive Data Types](#)

[Complex Data Types](#)

[Databases](#)

[Tables](#)

[Partitions](#)

[Buckets](#)

[Views](#)



# Partitions

To eliminate unnecessary scans of entire table when only a fragment is needed a table can be divided into **partitions**

A **partition** corresponds to predefined columns and it is stored as subfolder in **HDFS**

When a table is searched only required **partitions** are accessed

Creating a partitioned table

```
CREATE TABLE IF NOT EXISTS part(  
  P_PARTKEY DECIMAL(12),  
  P_NAME VARCHAR(55),  
  P_TYPE VARCHAR(25),  
  P_SIZE DECIMAL(12),  
  P_COMMENT VARCHAR(23) )  
PARTITIONED BY (P_BRAND VARCHAR(20))  
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '|'   
  STORED AS TEXTFILE;
```

# Partitions

A partition must be added before data is loaded

```
ALTER TABLE part ADD PARTITION (P_BRAND='GoldenBolts');
```

Adding a partition

```
show partitions part;  
OK  
p_brand=GoldenBolts  
Time taken: 0.072 seconds, Fetched: 1 row(s)
```

Listing partitions

A command that loads a file into a table can be used to load a partition

```
LOAD DATA LOCAL INPATH '/local/home/janusz/HIVE-EXAMPLES/TPCHR/part.txt'  
OVERWRITE INTO TABLE part PARTITION (P_BRAND='GoldenBolts');
```

Loading data into a partition

A partition is stored in HDFS as a subfolder

```
hadoop fs -ls /user/hive/warehouse/part  
Found 1 items  
drwxrwxr-x - janusz supergroup 0 2017-07-01  
19:00 /user/hive/warehouse/part/p_brand=GoldenBolts
```

Finding a partition in HDFS

# Hive Data Structures

## Outline

[Primitive Data Types](#)

[Complex Data Types](#)

[Databases](#)

[Tables](#)

[Partitions](#)

[Buckets](#)

[Views](#)

# Buckets

Another way to speed up processing of a table is to divide it into **buckets**

A **bucket** corresponds to segment of file in **HDFS**

The values in a **bucket column** will be hashed by a user defined number into buckets.

Creating a table with buckets

```
CREATE TABLE customer(  
  C_CUSTKEY DECIMAL(12),  
  C_NAME VARCHAR(25),  
  C_PHONE CHAR(15),  
  C_ACCTBAL DECIMAL(12,2) )  
  CLUSTERED BY (C_CUSTKEY) INTO 2 BUCKETS  
  ROW FORMAT DELIMITED FIELDS TERMINATED BY '|';
```

Setting MapReduce and Hive parameters

```
set map.reduce.tasks = 2;  
set hive.enforce.bucketing = true;
```

# Buckets

**INSERT** can be used to populate a bucket table

Inserting into a table with buckets

```
INSERT INTO customer
values(1, 'Customer#000000001', '25-989-741-2988', 711.56);
INSERT INTO customer
values(2, 'Customer#000000002', '23-768-687-3665', 121.65);
INSERT INTO customer
values(3, 'Customer#000000003', '11-719-748-3364', 7498.12);
INSERT INTO customer
values(4, 'Customer#000000004', '14-128-190-5944', 2866.83);
INSERT INTO customer
values(5, 'Customer#000000005', '13-750-942-6364', 794.47)
```

# Hive Data Structures

## Outline

[Primitive Data Types](#)

[Complex Data Types](#)

[Databases](#)

[Tables](#)

[Partitions](#)

[Buckets](#)

[Views](#)

# Views

**Views** are logical data structures that simplify queries

**Views** do not store data or get materialized

**Once a view** is created its definition is frozen and changes in the tables used in the view definition are not reflected in the view schema

```
CREATE VIEW vcustomer AS
  SELECT C_CUSTKEY, C_NAME, C_PHONE
  FROM CUSTOMER
  WHERE C_CUSTKEY < 5;
```

Creating a view

# References

Gross C., Gupta A., Shaw S., Vermeulen A. F., Kjerrumgaard D., Practical Hive: A guide to Hadoop's Data Warehouse System, Apress 2016, Chapter 4 (Available through UOW library)

Lee D., Instant Apache Hive essentials how-to: leverage your knowledge of SQL to easily write distributed data processing applications on Hadoop using Apache Hive, Packt Publishing Ltd. 2013 (Available through UOW library)

[Apache Hive](#)