


Pathfinding

Overview

- Graphs
- Pathfinding algorithms
 - Dijkstra
 - A*
- Pathfinding networks
- Advance techniques
 - Group path-finding
 - Hierarchical path-finding


The Pathfinding Graph

- A graph consists of
 - Nodes
 - Connections
 - An unordered pair of nodes
- Pathfinding algorithms use
 - Non-negative weighted graph
 - A path consists of
 - Zero or more connections





The Pathfinding Graph

- Weighted graphs
 - Connections have weights/costs
- Costs
 - In games often represent
 - Time or distance
 - Can also be other factors
 - Non-negative




The Pathfinding Graph

- Total path cost
 - Sum connection costs
- Nodes are representative points
 - Can be centre of a region
 - Or some other position




The Pathfinding Graph

- Directed weighted graphs
 - Connections and costs in one direction only
 - If can return, considered as two connections
 - Connections now ordered
 - E.g., from node 1 to node 2



The Pathfinding Graph

- Representation
 - Need to represent graph for pathfinding algorithms
 - Lookup table
 - Easy to implement but impractical for applications with a large number of nodes



	To	A	B	C	D	E	F	G	H
From	A	X	C	C	C	C	H	H	H
	B	C	X	C	C	C	C	C	C
	C	A	B	X	D	E	D	D	A
	D	C	C	C	X	E	F	F	F
	E	C	C	C	D	X	D	D	C
	F	G	D	D	D	D	X	G	G
	G	H	H	H	F	F	F	X	H
	H	A	A	A	G	A	G	G	X

The Pathfinding Graph

- Representation

- Algorithms need to find outgoing connections from any given node

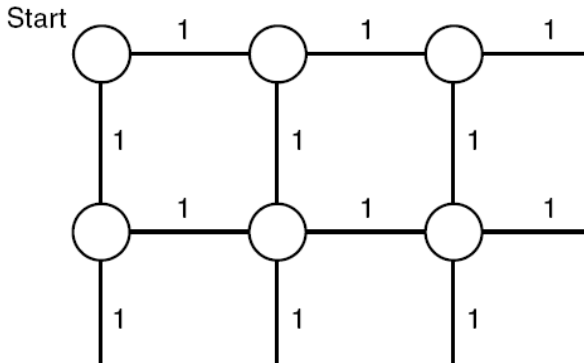
```
class Connection:
    // Returns the non-negative cost of the connection
    def getCost()


    // Returns the node that this connection came from
    def getFromNode()

    // Returns the node that this connection leads to
    def getToNode()

class Graph:
    // Returns an array of connections (of class Connection)
    // outgoing from the given node
    def getConnections(fromNode)
```


Dijkstra

- Dijkstra's algorithm
 - Designed to solve the 'shortest path' problem
 - Simpler version of the A* algorithm
 - Problem statement
 - Given a directed non-negative weighted graph, start and goal nodes, generate a path such that the total path cost is minimal among all possible paths
 - May be more than one with same minimal cost
 - If so, any one will do
- 
- ```
graph LR; S((Start)) ---|1| T1(()); S ---|1| T2(()); S ---|1| T3(()); T1 ---|1| T2; T1 ---|1| B1(()); T2 ---|1| T3; T2 ---|1| B2(()); T3 ---|1| B3(()); B1 ---|1| B2; B1 ---|1| B3; B2 ---|1| B3;
```



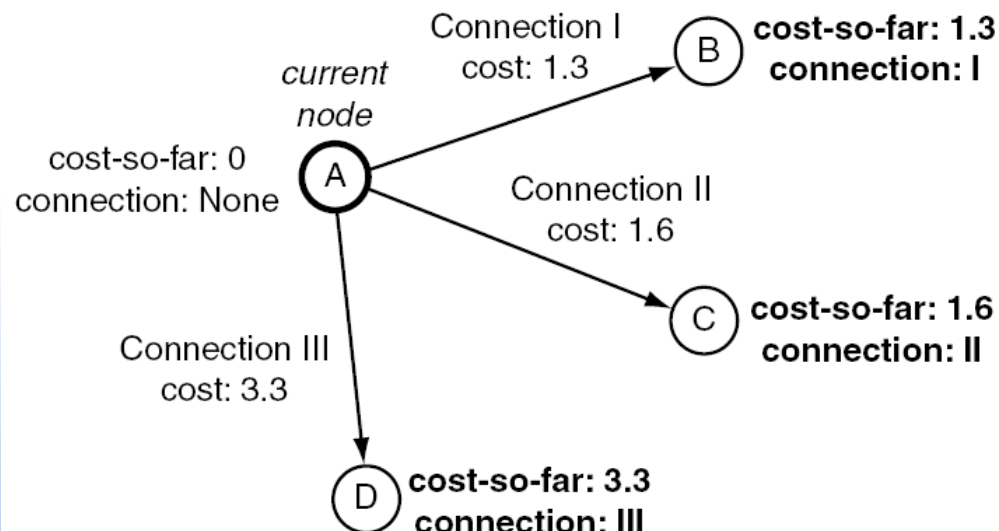
# Dijkstra

- In general, the path returned
  - Consists of a set of connections, not nodes
    - Two nodes may be linked by more than one connection
    - Each connection may have a different cost
- The algorithm
  - Works in iterations
    - Each iteration, considers one node (i.e., the “**current node**”) and follows outgoing connections
    - First iteration, considers the **start node**

# Dijkstra

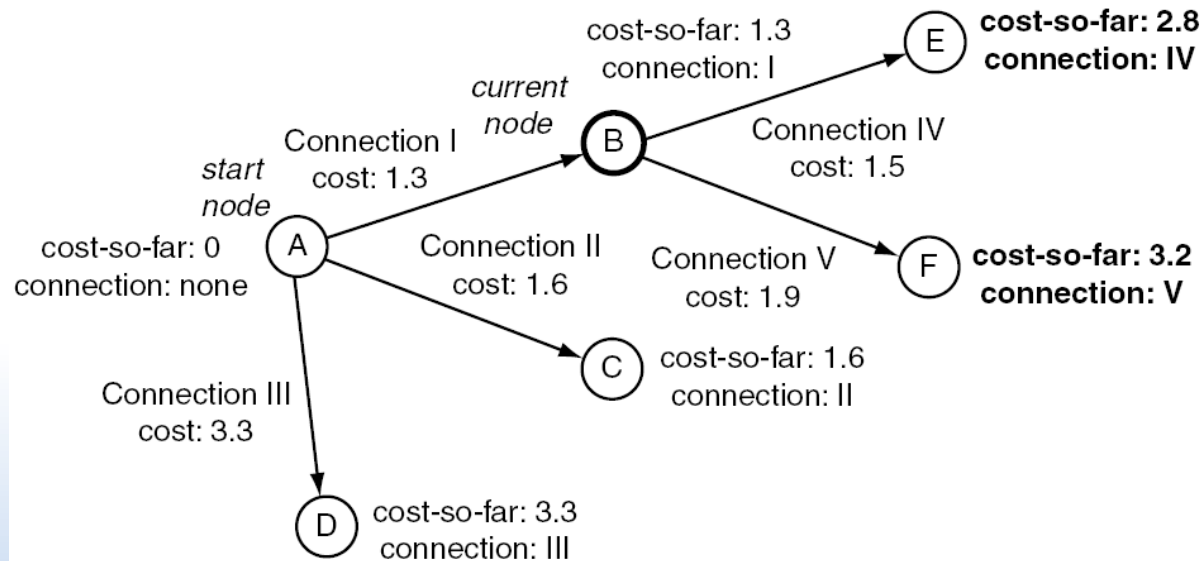
## ➤ Processing the current node

- Trace all the outgoing connections from the current node
- For each connection, find the end node and store
  1. Total cost of path thus far (**'cost-so-far'**), and
  2. The **connection** it arrived there from
- Need to record the cost-so-far and connection for each node



# Dijkstra

- For iterations after the first
  - Cost-so-far for the end node of each connection, is the
  - Sum of **connection cost** + **cost-so-far** of the current node



# Dijkstra

## ➤ Node lists

- Keeps track of all nodes seen so far in two lists, “open” list and “closed” list
- Can think of nodes in three categories
  - Closed: has been **processed in its own iteration**
  - Open: all nodes **seen but not processed**
  - Unvisited (in neither list)
- Each iteration
  - Choose a node from the open list with the **smallest cost-so-far**
    - » Set it as the current node
  - Process the current node
    - » Once processed, remove the current node from the open list and place it in the closed list

# Dijkstra

## ➤ Example

- At the start, only the start node will be on the open list
- For the first iteration, select the start node as the current node

*start  
node*

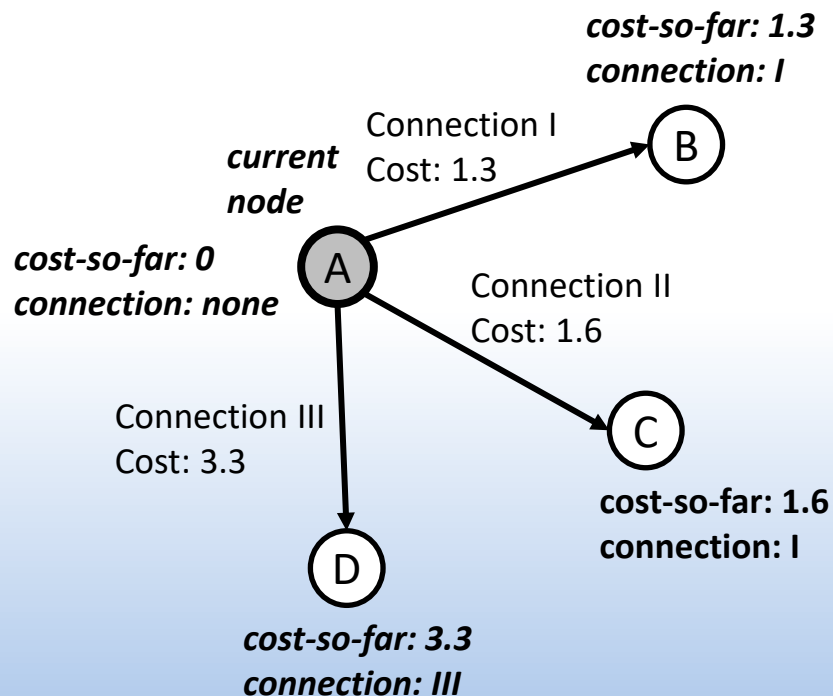
*cost-so-far: 0*  
*connection: none*

**A**

| Open List | Closed List |
|-----------|-------------|
| A         | -           |

# Dijkstra

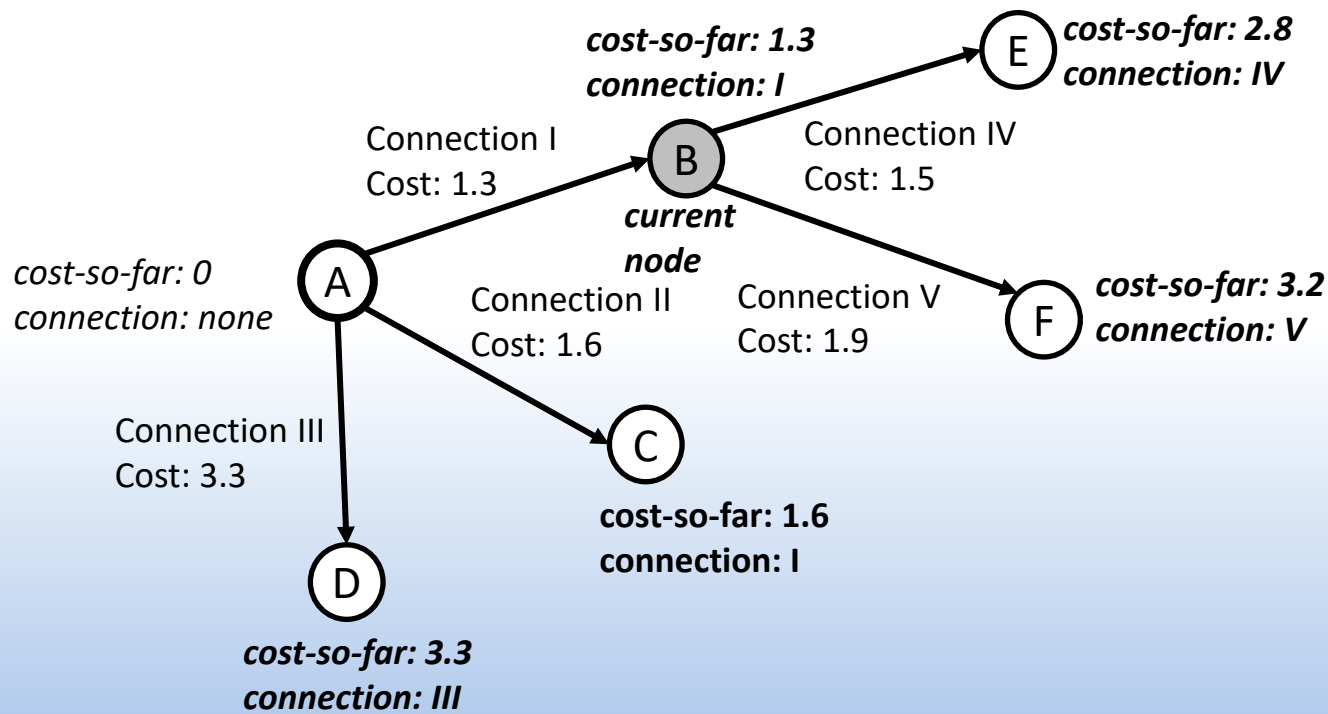
- Process the current node
  - Trace outgoing connections, find their end nodes, put these on the open list and update their node records
  - Once complete, remove the current node from the open list and put it in the closed list



| Open List | Closed List |
|-----------|-------------|
| B, C, D   | A           |

# Dijkstra

- Next iteration
  - Select the node with the lowest cost-so-far from the open list to the current node
  - Process the current node

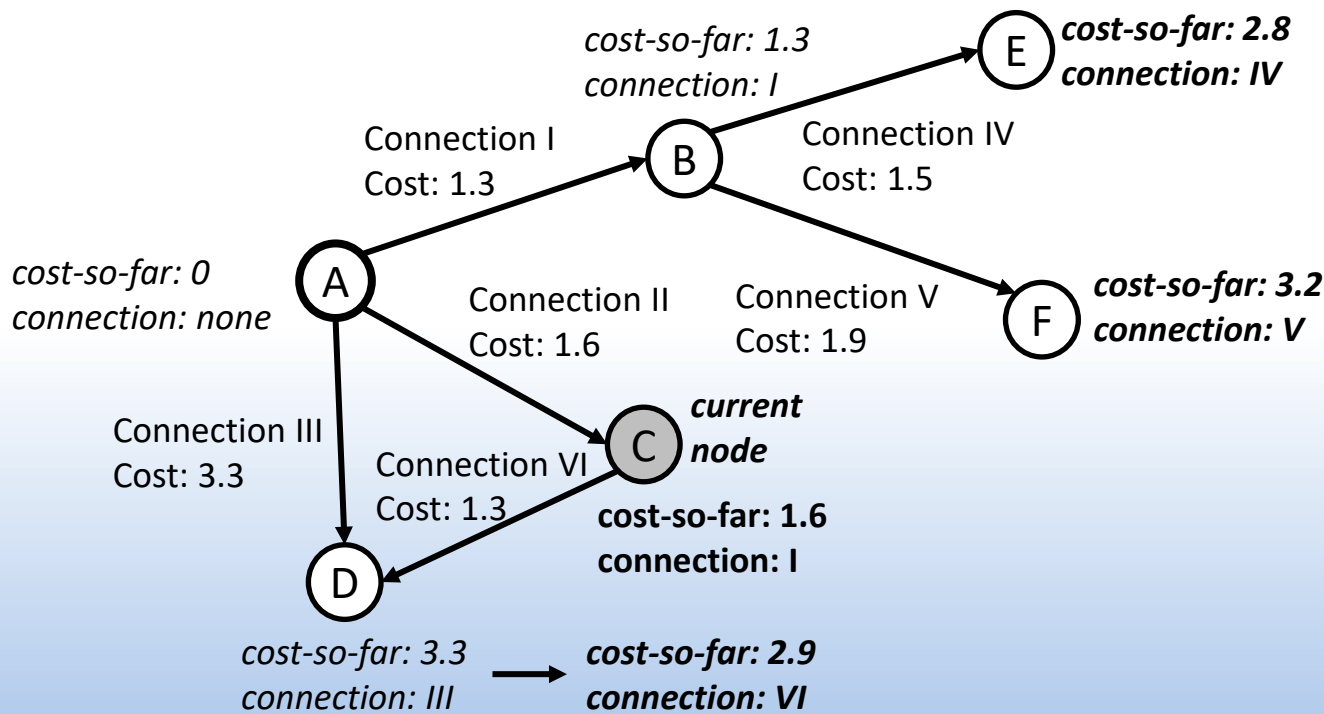


| Open List  | Closed List |
|------------|-------------|
| C, D, E, F | A, B        |



# Dijkstra

- Potential complication: end node already has a node record
  - Check for a better route
    - » If the cost-so-far of the current route is higher than the recorded value, don't do anything
    - » Otherwise, update the node records and put the node back onto the open list (if it was on the closed list)



| Open List | Closed List |
|-----------|-------------|
| D, E, F   | A, B, C     |

# Dijkstra

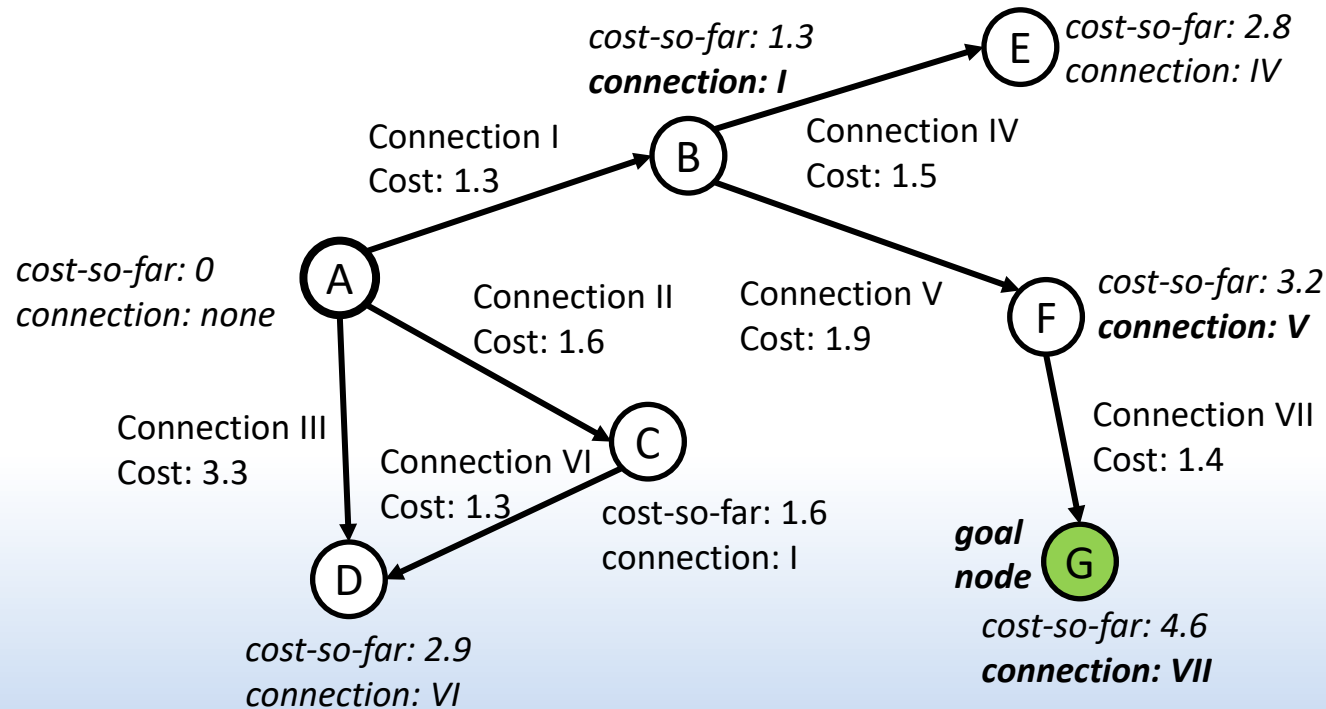
## ➤ Terminating the algorithm

- The basic algorithm terminates when the open list is empty
- In pathfinding
  - Only interested in reaching goal node, may stop earlier
  - Can terminate when the goal node has the smallest cost-so-far on the open list
- Why don't we stop as soon as the goal node is found?
  - No guarantee that this is the shortest path

# Dijkstra

## ➤ Retrieving the path

- Start at the goal node, work backward through the connections until the start node, then reverse the order




Connections working back from the goal: **VII, V, I**

Final path: **I, V, VII**

# Dijkstra

## ➤ Weaknesses

- Designed to search the whole graph indiscriminately
  - To work out shortest possible route
- Inefficient for 'point-to-point' path-finding
  - In general, want to consider as few nodes as possible
  - Number of nodes considered but not explored, called the 'fill' of the algorithm



# A\* Pathfinding

- A\* algorithm
  - Designed for point-to-point path-finding and not to solve shortest path problem
  - Works very similar to Dijkstra
    - But rather than always considering node on open list with the lowest cost-so-far
    - Choose the node 'most likely' to lead to the shortest overall path
    - Notion of 'most likely' controlled by a **heuristic**
      - If heuristic accurate, efficient
      - If heuristic terrible, can perform worse than Dijkstra
      - If heuristic is all 0, then exactly the same as Dijkstra

# A\* Pathfinding

- The algorithm
  - Like Dijkstra, works in iterations
  - Processing current node
    - Difference from Dijkstra? Store another value called the **“estimated-total-cost”**
      - Estimate of total cost for a path from the start node, through the current node, onto the goal
      - Estimated-total-cost = **cost-so-far + heuristic**
      - The heuristic is an estimate of the cost from that node to the goal
        - » Cannot be negative
      - This estimate is not part of the algorithm and is generated by a separate piece of code

# A\* Pathfinding

## ➤ Example

- At the start, only the start node will be on the open list
- For the first iteration, select the start node as the current node

*heuristic: 4.2*

*cost-so-far: 0*

*connection: none*

*estimated-total-cost: 4.2*

*start  
node*



| Open List | Closed List |
|-----------|-------------|
| A         | -           |

# A\* Pathfinding

- Process the current node
  - Trace outgoing connections, find their end nodes, put these on the open list and update their node records
  - Once complete, remove the current node from the open list and put it in the closed list

*heuristic: 4.2*

*cost-so-far: 0*

*connection: none*

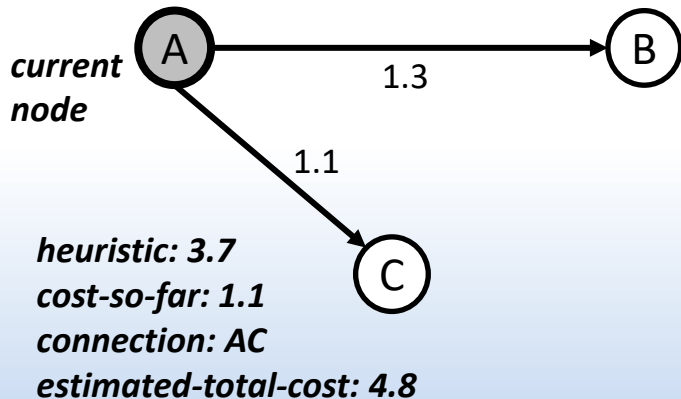
*estimated-total-cost: 4.2*

*heuristic: 3.2*

*cost-so-far: 1.3*

*connection: AB*

*estimated-total-cost: 4.5*

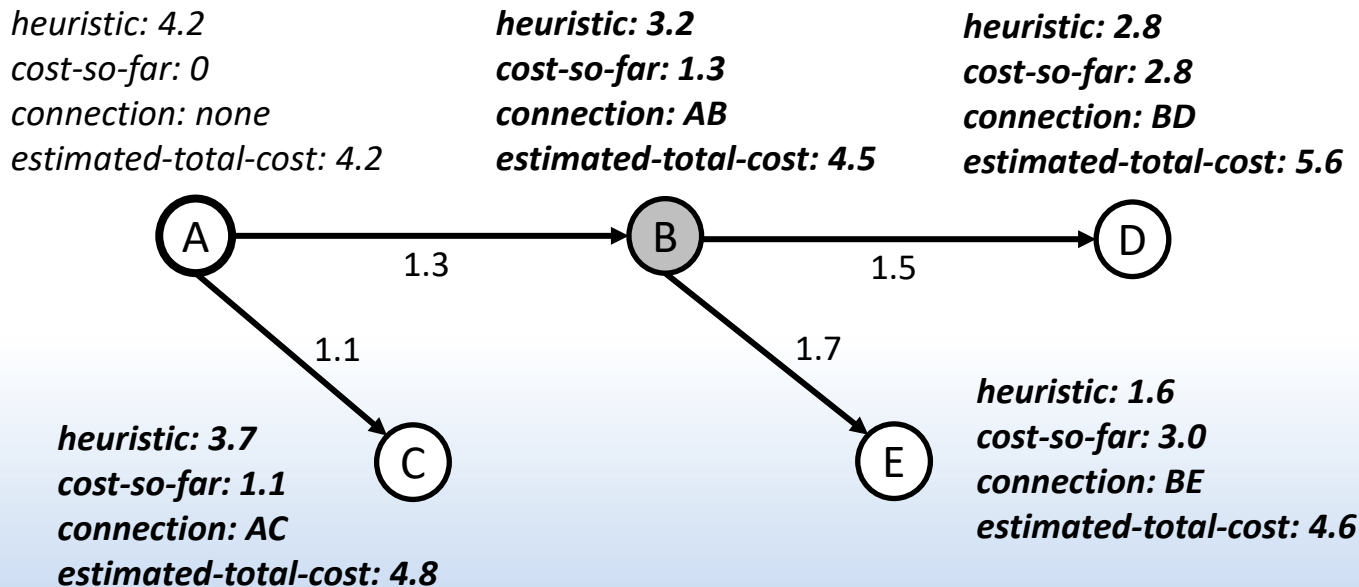


| Open List | Closed List |
|-----------|-------------|
| B, C      | A           |



# A\* Pathfinding

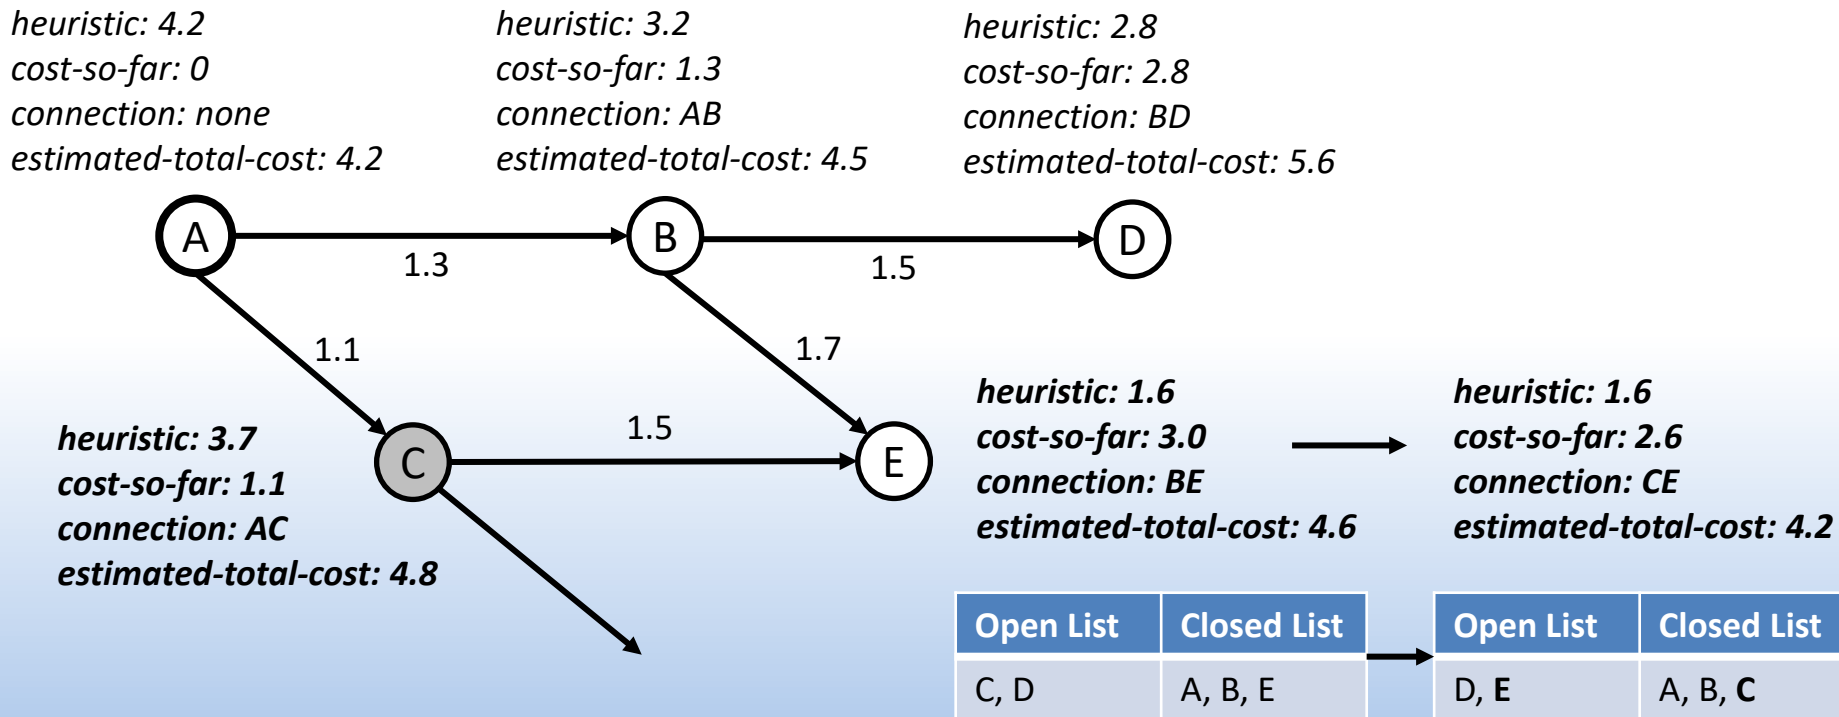
- Difference with Dijkstra
  - For the current node, select the node on the open list with **smallest estimated-total-cost**
  - **NOT** smallest cost-so-far



| Open List | Closed List |
|-----------|-------------|
| C, D, E   | A, B        |

# A\* Pathfinding

- Complication and solution
  - Compare the cost-so-far (the only reliable value)
    - » NOT the estimated-total-cost (only and estimated cost)
  - If node was on the closed list, put it back on the open list



# A\* Pathfinding

## ➤ Terminating the algorithm

- When the goal node has the smallest value on the open list
- A\* completely relies on the fact that can theoretically produce non-optimal results
  - Depends on the choice of heuristic function
  - Many implementations just terminate upon finding the goal

## ➤ Retrieving the path

- Start a goal node, work back until the start node (same as Dijkstra)

# A\* Pathfinding


- Heuristic function
  - Inconvenient to produce a different heuristic function for each possible goal
    - Often parameterised by the goal node
  - Might involve some algorithmic process
    - If the process is complex, time spent evaluating heuristics might dominate the pathfinding algorithm
  - For non-perfect heuristics
    - A\* behaves slightly differently depending on whether underestimating or overestimating

# A\* Pathfinding

- Common heuristics in games

- Euclidean distance

- If connection cost represents distance
    - Distance “as the crow flies”, measured between two points through obstructions
    - Either accurate or guaranteed to underestimate
    - In outdoor settings with few constraints, can be very fast and accurate



# A\* Pathfinding

## ➤ Euclidean distance

- Dramatically underestimates in indoor settings
  - Might cause less than optimal pathfinding

Indoor level

```
x xxx xxx oooooo. . .
x xxx xxx Oxxxxx .
x x x xxxxxx ooo
x x xxxxx xxxxxx xxx
xxx xxxxxxxx xxxxxx xxx
x xxxx x xxx
x . . . xxx xxxxxx x
x . . . x xxxxxx x
x . . . xxx xxxxxxxxxx
x . xxxxxx xxxxxx x
x . xxxxxx xxxxxx xxx
x o xxxxxx xxxxxx x
x x xxxxxx xxxxxx xx
x x
x xxxxxxxx . Oxx x
x xxxxxxxx . Oxx xx
x xxxxxxxx . Oxxxxxx
x xxxxxxxx . Oxxx
x xxxxxxxx Oxx
x xxxxxxxx x
xxxxxxx Oxxxxxxx
```

Outdoor level

```
xxoo.
xxxoo.
Oxxxxoo.
Oxxxxoo.
. Oxxxxoo.
. . Oxxxxoo.
. . . Oxxxxoo.
. . . . Oxx xoo.
. oo Oxxoo.
. Oxxoo.
. Oxxoo.
. Oxxoo. . . .
. Oxxoo. . .
. Oxxoo.
. Oxxo
. Oxx
. O
```


### Key

- × Closed node
- o Open node
- Unvisited node

# A\* Pathfinding

## ➤ Cluster heuristic

- Group nodes together in clusters
- Nodes in clusters represent highly interconnected region
- Often done manually
- Lookup table gives smallest path length



# A\* Pathfinding


- Often dramatically improves performance in indoor areas
  - As it takes convoluted routes into account
- However, all nodes in a cluster give the same heuristic value
  - A\* cannot easily find the best route through a cluster
- A cluster tends to be almost completely filled before the algorithm moves to the next cluster
- Tradeoff
  - If a cluster is small, the accuracy can be excellent
    - » But the lookup table can be large
  - If too large, marginal performance gain
    - » Simpler heuristic better



# A\* Pathfinding

## ➤ Cluster heuristic

- Fill visualisation for an indoor setting



# A\* Pathfinding

## ► Cluster heuristic

- Fill visualisation for an outdoor setting

## Euclidean distance heuristic

[illegible]

Null heuristic

## Key



- × Closed node
- Open node
- Unvisited node

# Pathfinding Networks

- Pathfinding network
  - Search a network of connected nodes
  - Network creation can be time-consuming
    - Generated before path-finding operations
  - Can be done manually, or generated automatically, or a mixture of both
  - Most games only use 2D network and most of these networks do not change during the game

# Pathfinding Networks

- Rectangular grids (also known as tiles)
  - Can be square or hexagonal
  - Can be used to automatically generate path-finding networks
  - Easy to generate algorithmically and determine connectivity
  - Easy to implement
    - Can use a simple 2D array





# Pathfinding Networks

- Hex grids
  - Slightly harder to implement
  - Advantage
    - » Six adjacent nodes all equidistant from the current node centre
    - » Movement tends to be more regular





# Pathfinding Networks

## ➤ Arbitrary grids

- Nodes only placed at important locations
- Usually done manually
- Advantages
  - Nodes can be placed exactly along the desired movement for path following
  - Needs to store less data
- Disadvantage
  - Manual work required

## ➤ Navigation mesh

- Divides the environment into a mesh, assign nodes to the mesh





# Advance Techniques

- Group pathfinding
  - Large number of pathfinding operations if done for each character
  - If all characters directed to move to the same location
    - Can use single path-finding operation
    - Then, share and/or adjust results
  - The simplest way is to find a path from the centre of group to the destination

# Advance Techniques

## ➤ Problem


- Units will all line up in a row (called 'antlining')
  - Takes long time for last unit to arrive
  - If enemy near destination, group easily destroyed
- To address antlining
  - Offset the individual unit paths
  - Formation much tighter while moving and at destination





# Advance Techniques

- Must make special allowance for width of group
  - Single unit pathfinding locates a path that is as wide as a unit
  - For a group, need to locate a wider path
  - Care must be taken in width checking
    - Might be better to take a narrow opening rather than go around the long way
    - In this case, antlining is perfectly acceptable




# Advance Techniques

## ➤ Antlining and regrouping




# Advance Techniques

- Hierarchical path-finding
  - Plan overview of route first then refine as needed



Original network contains 56 nodes and high level of detail




Hierarchical network contains less detailed representation and fewer nodes (12 nodes)



# Advance Techniques

- Intermediate destinations
  - Creating shorter paths




Path from A to Z

- Search on hierarchical network  $H1 \rightarrow H6 \rightarrow H7 \rightarrow H8 \rightarrow H12$
- Shorter path from A to Z:  $A \rightarrow H6 \rightarrow H7 \rightarrow H8 \rightarrow Z$

# Advance Techniques

## ➤ Pruning the network

- Remove unrelated nodes from consideration



## Path from A to Z

- Search on hierarchical network  $H1 \rightarrow H6 \rightarrow H7 \rightarrow H8 \rightarrow H12$
- After pruning only need to search 24 rather than full 56 nodes

# References

- Among others, material sourced from
  - <https://docs.unity3d.com>
  - Jason Gregory, Game Engine Architecture, A.K. Peters
  - Ian Millington, Artificial Intelligence for Games, Morgan Kaufmann
  - Mat Buckland, Programming Game AI by Example, Wordware Publishing
  - John Alquist and Jeannie Novak, Game Development Essentials: Game Artificial Intelligence, Delmar Publishing
  - Screenshots are from various games, mainly sourced from [www.ign.com](http://www.ign.com) and [www.mobygames.com](http://www.mobygames.com)